Notation: in this note, y is a scale random variable, and x is a $K \times 1$ random vector.

1 Conditional Expectation

We view the regression as a problem of supervised learning. Supervised learning uses a function of x, say, $g(x)$, to predict y. We do not x can perfectly predict y; otherwise their relationship is deterministic. The prediction error

$$e = y - g(x)$$

depends on the choice of g. There are numerous possible choices of g. Which one is the best? To answer this question, we need to decide a criterion to compare different g, which is called the *loss function* $L(y, g(x))$. A particularly convenient one is the *quadratic loss*, defined as

$$L(y, g(x)) = (y - g(x))^2.$$

Since the data is random, we average the loss function across the joint distribution of (y, x) as $R(y, g(x)) = E[L(y, g(x))]$, which is called the *risk*. For the quadratic loss function, the corresponding risk

$$R(y, g(x)) = E\left[(y - g(x))^2\right],$$

is called the *mean squared error* (MSE). MSE is a deterministic quantity since the randomness is integrated out.

What is the optimal choice of g if we aim to minimize the MSE?

Proposition 1. *The conditional mean function $m(x) = E[y|x]$ minimizes MSE.*

Before we prove the above proposition, we first discuss some properties of the conditional mean function. Obviously

$$y = m(x) + (y - m(x)) = m(x) + \epsilon,$$

where $\epsilon \equiv y - m(x)$ is called the *regression error*. This equation holds for (y, x) following any joint distribution, as long as $E[y|x]$ exists. The error term ϵ satisfies these properties:

- $E[\epsilon|x] = E[y - m(x)|x] = E(y|x) - m(x) = 0$,
- $E[\epsilon] = E[E[\epsilon|x]] = E(0) = 0$,
- For any function $h(x)$, we have $E[h(x)\epsilon] = E[E[h(x)\epsilon|x]] = E[h(x) \cdot E[\epsilon|x]] = 0$.

The last property implies that ϵ is uncorrelated with any function of x. In particular, when h is the identity function $h(x) = x$, we have $E[x\epsilon] = \operatorname{cov}(x, \epsilon) = 0$.

Proof of Proposition 1. The optimality of the conditional mean can be confirmed by "guess-and-verify." For an arbitrary $g(x)$, the MSE can be decomposed into three terms

$$\begin{aligned} E\left[(y - g(x))^2\right] &= E\left[(y - m(x) + m(x) - g(x))^2\right] \\ &= E\left[(y - m(x))^2\right] + 2E\left[(y - m(x))(m(x) - g(x))\right] + E\left[(m(x) - g(x))^2\right]. \end{aligned}$$

.The first term is irrelevant to $g(x)$. The second term

$$\begin{aligned} 2E\left[\left(y-m\left(x\right)\right)\left(m\left(x\right)-g\left(x\right)\right)\right] &= 2E\left[\epsilon\left(m\left(x\right)-g\left(x\right)\right)\right] \\ &= 2E\left[E\left[\epsilon\left(m\left(x\right)-g\left(x\right)\right)|x\right]\right] \\ &= 2E\left[\left(m\left(x\right)-g\left(x\right)\right)E\left[\epsilon|x\right]\right]=0. \end{aligned}$$

is again irrelevant of $g(x)$. The third term, obviously, is minimized at $g(x)=m(x)$. □

Our perspective so far deviates from mainstream econometric textbooks, most of which start the regression model by assuming that the dependent variable y is generated by, or is modeled as, $g(x)+\epsilon$ for some unknown function $g(\cdot)$ and some error term ϵ such that $E[\epsilon|x]=0$. Instead, we take a predictive framework regardless the data generating process. What we observe are y and x and we are solely interested in seeking a function $g(x)$ to predict y as accurately as possible under the MSE criterion.

2 Linear Projection

As discussed in the previous section, the conditional mean function $m(x)$ is the function that minimizes the MSE. However,

$$m\left(x\right)=E\left[y|x\right]=\int yf\left(y|x\right)\mathrm{d}y$$

is a complex function of x, as it depends on the joint distribution of (y,x), which is almost always unknown in practice. Now let us make the prediction task even simpler. How about we minimize the MSE within all linear functions in the form of $g(x)=x'b$ for $\beta\in\mathbb{R}^K$? The minimization problem is

$$\min_{b\in\mathbb{R}^K}E\left[\left(y-x'b\right)^2\right]. \tag{1}$$

Take the first-order condition of the MSE

$$\frac{\partial}{\partial b}E\left[\left(y-x'b\right)^2\right]=-2E\left[x\left(y-x'b\right)\right]=0.$$

Rearrange the above equation and we solve the optimal b as

$$\beta=\left(E\left[xx'\right]\right)^{-1}E\left[xy\right]$$

if $E[xx']$ is invertible. The function $x'\beta$ is called the *best linear projection* of y on x, and the vector β is called the *linear projection coefficient*.

Remark 2. The linear function is not as restrictive as one might thought. It can be used to produce some nonlinear (in random variables) effect if we re-define x. For example, if

$$y=x_1\beta_2+x_2\beta_2+x_1^2\beta_3+e,$$

then $\frac{\partial}{\partial x_1}m\left(x_1,x_2\right)=\beta_1+x_1\beta_3$, which is nonlinear in x_1, while it is still linear in the parameter β if we define a set of new regressors as $(\tilde{x}_1,\tilde{x}_2,\tilde{x}_3)=\left(x_1,x_2,x_1^2\right)$.

Remark 3. If (y, x) is jointly normal in the form $\begin{pmatrix} y \\ x \end{pmatrix} \sim \mathrm{N}\left(\begin{pmatrix} \mu_y \\ \mu_x \end{pmatrix}, \begin{pmatrix} \sigma_y^2 & \rho\sigma_y\sigma_x \\ \rho\sigma_y\sigma_x & \sigma_x^2 \end{pmatrix} \right)$ where ρ is the correlation coefficient, then

$$E\left[y|x\right] = \mu_y + \rho\frac{\sigma_y}{\sigma_x}\left(x - \mu_x\right) = \left(\mu_y - \rho\frac{\sigma_y}{\sigma_x}\mu_x\right) + \rho\frac{\sigma_y}{\sigma_x}x,$$

is a liner function of x. In this example, the conditional mean coincides with a linear function.

Remark 4. Even though in general $m\left(x\right) \neq x'\beta$, the linear form $x'\beta$ is still useful in approximating $m\left(x\right)$. That is, $\beta = \arg\min_{b\in\mathbb{R}^K} E\left[\left(m(x) - x'b\right)^2\right]$.

Proof. $\arg\min_{b\in\mathbb{R}^K} E\left[\left(m(x) - x'b\right)^2\right]$

First order condition gives $\frac{\partial}{\partial b}E\left[\left(m(x) - x'b\right)^2\right] = -2xE[(m(x) - x'b)] = 0$. Rearrange the terms and obtain

$$\begin{aligned} E[x \cdot m(x)] &= E(xx')\beta \\ \beta &= [E(xx')]^{-1}E[E[xy|x]] \\ &= [E(xx')]^{-1}E(xy). \end{aligned}$$

Thus β is also the best linear approximation to $m\left(x\right)$ under MSE. □

We may rewrite the linear regression model, or the *linear projection model,* as

$$\begin{aligned} y &= x'\beta + e \\ E[xe] &= 0, \end{aligned}$$

where $e = y - x'\beta$ is called the *projection error*, to be distinguished from $\epsilon = y - m(x)$.
Exercise: show (a) $E\left[xe\right] = 0$. (b) If x contains a constant, then $E\left[e\right] = 0$.

2.1 Omitted Variable Bias

We write the *long regression* as

$$y = x_1'\beta_1 + x_2'\beta_2 + \beta_3 + \varepsilon,$$

and the *short regression* as

$$y = x_1'\gamma_1 + \gamma_2 + u.$$

If β_1 in the long regression is the parameter of interest, omitting x_2 as in the short regression will render *omitted variable bias* (meaning $\gamma_1 \neq \beta_1$) unless x_1 and x_2 are uncorrelated.

We first demean all the variables in the two regressions, which is equivalent as if we project out the effect of the constant. The long regression becomes

$$\tilde{y} = \tilde{x}_1'\beta_1 + \tilde{x}_2'\beta_2 + \tilde{\varepsilon},$$

and the short regression becomes

$$\tilde{y} = \tilde{x}_1'\gamma_1 + \tilde{u},$$

where *tilde* denotes the demeaned variable.

After demeaning, the cross-moment equals to the covariance. The short regression coefficient

$$\begin{aligned}
\gamma_1 &= \left(E\left[\tilde{x}_1\tilde{x}_1'\right]\right)^{-1} E\left[\tilde{x}_1\tilde{y}\right] \\
&= \left(E\left[\tilde{x}_1\tilde{x}_1'\right]\right)^{-1} E\left[\tilde{x}_1\left(\tilde{x}_1'\beta_1 + \tilde{x}_2'\beta_2 + e\right)\right] \\
&= \left(E\left[\tilde{x}_1\tilde{x}_1'\right]\right)^{-1} E\left[\tilde{x}_1\tilde{x}_1'\right]\beta + \left(E\left[\tilde{x}_1\tilde{x}_1'\right]\right)^{-1} E\left[\tilde{x}_1\tilde{x}_2'\right]\beta_2 \\
&= \beta_1 + \left(E\left[\tilde{x}_1\tilde{x}_1'\right]\right)^{-1} E\left[\tilde{x}_1\tilde{x}_2'\right]\beta_2.
\end{aligned}$$

Therefore, $\gamma_1 = \beta_1$ if and only if $E\left[\tilde{x}_1\tilde{x}_2'\right]\beta_2 = 0$, which demands either $E\left[\tilde{x}_1\tilde{x}_2'\right] = 0$ or $\beta_2 = 0$.

Obviously we prefer to run the long regression to attain β_1 if possible, as it is a model general model than the short regression. However, sometimes x_2 is simply unobservable so the long regression is infeasible. When only the short regression is available, in some cases we are able to sign the bias, meaning that we know whether γ_1 is bigger or smaller than β_1.